

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 01

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
data=pd.read_csv('/content/Iris_Dataset.csv')
data
```

```
   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  variety
0  1  5.1  3.5  1.4  0.2  Iris-setosa
1  2  4.9  3.0  1.4  0.2  Iris-setosa
2  3  4.7  3.2  1.3  0.2  Iris-setosa
3  4  4.6  3.1  1.5  0.2  Iris-setosa
4  5  5.0  3.6  1.4  0.2  Iris-setosa
... ..
145 146 6.7 3.0 5.2 2.3  Iris-virginica
146 147 6.3 2.5 5.0 1.9  Iris-virginica
147 148 6.5 3.0 5.2 2.0  Iris-virginica
148 149 6.2 3.4 5.4 2.3  Iris-virginica
149 150 5.9 3.0 5.1 1.8  Iris-virginica
150 rows x 6 columns
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 # Column Non-Null Count Dtype
---  ---
0 Id 150 non-null int64
1 SepalLengthCm 150 non-null float64
2 SepalWidthCm 150 non-null float64
3 PetalLengthCm 150 non-null float64
4 PetalWidthCm 150 non-null float64
5 variety 150 non-null object
```

```
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```


```
data.describe()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150 000000	7 900000	4 400000	6 900000	2 500000

```
data.value_counts('variety')
```

	count
variety	
Iris-setosa	50
Iris-versicolor	50
Iris-virginica	50

```
sns.countplot(x='variety',data=data,)
plt.show()
```



```
dummies=pd.get_dummies(data.variety)
```


```
FinalDataset=pd.concat([pd.get_dummies(data.variety),data.iloc[:,[0,1,2,3]]],
axis=1)
```

```
FinalDataset.head()
```


	Iris-setosa	Iris-versicolor	Iris-virginica	Id										
SepalLengthCm	SepalWidthCm	PetalLengthCm	0	True										
False	False	1	5.1	3.5	1.4	1	True	False	False	2	4.9	3.0	1.4	2
True	False	False	3	4.7	3.2	1.3	3	True	False	False	4	4.6	3.1	
1.5	4	True	False	False	5	5	0	3	6	1	4			

```
sns.scatterplot(x='SepalLengthCm',y='SepalWidthCm',hue='variety',data=data,)
```

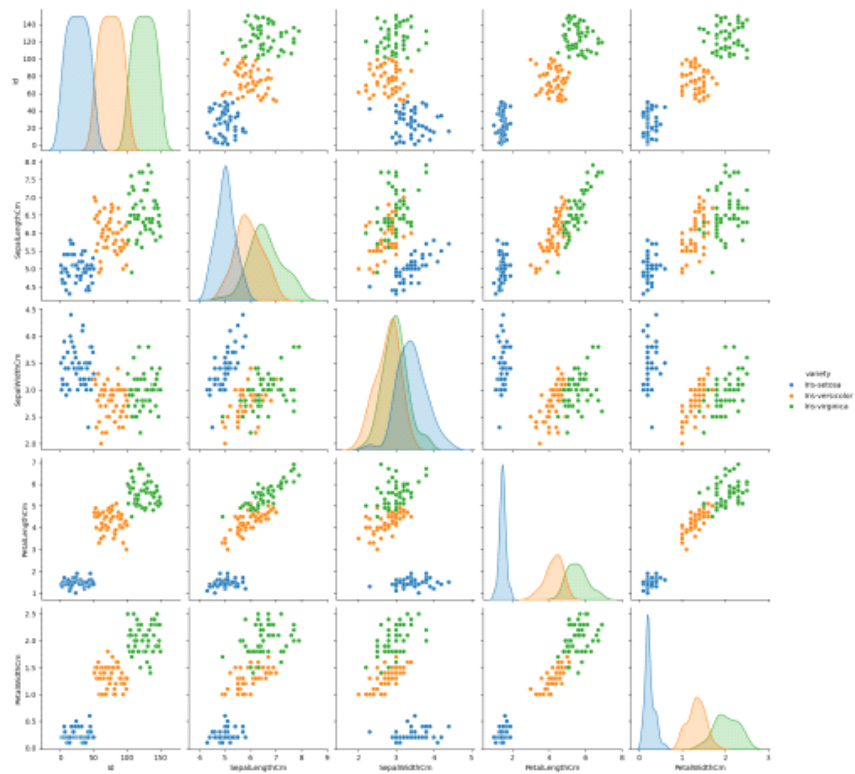
```
<Axes: xlabel='SepalLengthCm', ylabel='SepalWidthCm'>
```




```
sns.scatterplot(x='PetalLengthCm',y='PetalWidthCm',hue='variety',data=data,)
```




```
sns.pairplot(data,hue='variety',height=3);
```




```
plt.show()
sns.FacetGrid(data,hue='variety',height=5).map(
sns.histplot,'PetalLengthCm').add_legend();
plt.show();
```



```
sns.FacetGrid(data,hue='variety',height=5).map(  
sns.histplot,'Petal.WidthCm').add_legend();  
plt.show();
```



```
sns.FacetGrid(data,hue='variety',height=5).map(
sns.histplot,'SepalLengthCm').add_legend();
plt.show();
```



```
sns.FacetGrid(data,hue='variety',height=5).map(sns.histplot,'SepalWidthCm').a
dd_legend();
plt.show();
```

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 02

```
import numpy as np
array=np.random.randint(1,100,9)
array

array([83, 25, 19, 47, 62, 15, 96, 39, 51])

np.sqrt(array)

array([9.11043358, 5. , 4.35889894, 6.8556546 , 7.87400787,
 3.87298335, 9.79795897, 6.244998 , 7.14142843])
```

```
array.ndim
```

```
1
```

```
new_array=array.reshape(3,3)
```

```
new_array
```

```
array([[83, 25, 19],  
       [47, 62, 15],  
       [96, 39, 51]])
```

```
new_array.ndim
```

```
2
```

```
new_array.ravel()
```

```
array([83, 25, 19, 47, 62, 15, 96, 39, 51])
```

```
newm=new_array.reshape(3,3)
```

```
newm
```

```
array([[83, 25, 19],  
       [47, 62, 15],  
       [96, 39, 51]])
```

```
newm[2,1:3]
```

```
array([39, 51])
```

```
newm[1:2,1:3]
```

```
array([[62, 15]])
```

```
new_array[0:3,0:0]
```

```
array([], shape=(3, 0), dtype=int64)
```

```
new_array[0:2,0:1]
```

```
array([[83],  
       [47]])
```

```
new_array[0:3,0:1]
```

```
array([[83],  
       [47],  
       [96]])
```

```
new_array[1:3]
```



```
array([[47, 62, 15],
       [96, 39, 51]])
```

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 03

```
import numpy as np
import pandas as pd
list=[[1,'Smith',50000],[2,'Jones',60000]]
```

```
df=pd.DataFrame(list)
df
```

```
   0 1 2
0 1 Smith 50000
1 2 Jones 60000
```

```
df.columns=['Empd','Name','Salary']
df
```

```
   Empd Name Salary
0 1 Smith 50000
1 2 Jones 60000
```

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2 entries, 0 to 1
Data columns (total 3 columns):
 # Column Non-Null Count Dtype
---  ---
0 Empd 2 non-null int64
1 Name 2 non-null object
2 Salary 2 non-null int64
dtypes: int64(2), object(1)
memory usage: 176.0+ bytes
```

```
df=pd.read_csv("/content/50_Startups.csv")
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
 # Column Non-Null Count Dtype
---  ---
0 R&D Spend 50 non-null float64
1 Administration 50 non-null float64
2 Marketing Spend 50 non-null float64
3 State 50 non-null object
4 Profit 50 non-null float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

```
df.head()
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
df.tail()
```

	R&D Spend	Administration	Marketing Spend	State	Profit
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41
49	0.00	116983.80	45173.06	California	14681.40

```
import numpy as np
import pandas as pd
df=pd.read_csv("/content/employee.csv")
```

```
df.head()
```

	emp id	name	salary
0	1	SREE VARSSINI K S	5000
1	2	SREEMATHI B	6000
2	3	SREYA G	7000

3 4 SREYASKARI MULLAPUDI 5000

4 5 SRI AKASH U G 8000

```
df.tail()
```

```
   emp id name salary
2 3 SREYA G 7000
3 4 SREYASKARI MULLAPUDI 5000
4 5 SRI AKASH U G 8000
5 6 SRI HARSHAVARDHANAN R 3000
6 7 SRI HARSHAVARDHANAN R 6000
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7 entries, 0 to 6
Data columns (total 3 columns):
 # Column Non-Null Count Dtype
---  ---
0 emp id 7 non-null int64
1 name 7 non-null object
2 salary 7 non-null int64
dtypes: int64(2), object(1)
memory usage: 296.0+ bytes
```

```
df.salary
```

```
   salary
0 5000
1 6000
2 7000
3 5000
4 8000
5 3000
6 6000
```

```
type(df.salary)
```

```
pandas.core.series.Series
def __init__(data=None, index=None, dtype: Dtype | None=None, name=None,
copy: bool | None=None,
fastpath: bool=False) -> None
```

One-dimensional ndarray with axis labels (including time series).
Labels need not be unique but must be a hashable type. The object supports both integer- and label-based indexing and provides a host of methods for performing operations involving the index. Statistical

```
th d f d h b idd t t ti ll l d

df.salary.mean()
5714.285714285715
df.salary.median()
6000.0

df.salary.mode()

  salary
0 5000
1 6000

df.salary.var()
2571428.5714285714

df.salary.std()
1603.5674514745463

df.describe()

   emp id salary
count 7.000000 7.000000
mean 4.000000 5714.285714
std 2.160247 1603.567451
min 1.000000 3000.000000
25% 2.500000 5000.000000
50% 4.000000 6000.000000
75% 5.500000 6500.000000
max 7 000000 8000 000000

df.describe(include='all')

   emp id name salary
count 7.000000 7 7.000000
unique NaN 6 NaN
```

top NaN SRI HARSHAVARDHANAN R NaN
freq NaN 2 NaN
mean 4.000000 NaN 5714.285714
std 2.160247 NaN 1603.567451
min 1.000000 NaN 3000.000000
25% 2.500000 NaN 5000.000000
50% 4.000000 NaN 6000.000000
75% 5.500000 NaN 6500.000000
max 7 000000 NaN 8000 000000

```
empCol=df.columns
```

```
empCol
```

```
Index(['emp id', 'name ', 'salary'], dtype='object')
```

```
emparray=df.values
```

```
emparray
```

```
array([[1, 'SREE VARSSINI K S', 5000],  
       [2, 'SREEMATHI B', 6000],  
       [3, 'SREYA G', 7000],  
       [4, 'SREYASKARI MULLAPUDI', 5000],  
       [5, 'SRI AKASH U G', 8000],  
       [6, 'SRI HARSHAVARDHANAN R', 3000],  
       [7, 'SRI HARSHAVARDHANAN R', 6000]], dtype=object)
```

```
employee_DF=pd.DataFrame(emparray,columns=empCol)
```

```
employee_DF
```

	emp id	name	salary
0	1	SREE VARSSINI K S	5000
1	2	SREEMATHI B	6000
2	3	SREYA G	7000
3	4	SREYASKARI MULLAPUDI	5000
4	5	SRI AKASH U G	8000
5	6	SRI HARSHAVARDHANAN R	3000
6	7	SRI HARSHAVARDHANAN R	6000

Lab experiments**Roll no: 230701017****Name: Ajay Srinivas R****Class: CSE-A II****Subject: Fundamentals of data science (CS23334)****Experiment: 04**

```
#sample calculation for low range(lr) , upper range (ur),percentile
import numpy as np
array=np.random.randint(1,100,16) # randomly generate 16 numbers between 1 to
100
array

    array([27, 50, 44, 6, 58, 61, 23, 86, 67, 20, 75, 7, 79, 61, 90, 54])

array.mean()

    50.5

np.percentile(array,25)

    26.0

np.percentile(array,50)

    56.0

np.percentile(array,75)

    69.0

np.percentile(array,100)

    90.0

#outliers detection
def outDetection(array):
    sorted(array)
    Q1,Q3=np.percentile(array,[25,75])
    IQR=Q3-Q1
    lr=Q1-(1.5*IQR)
    ur=Q3+(1.5*IQR)
    return lr,ur


lr,ur=outDetection(array)

lr,ur
```

(-38.5, 133.5)


```
import seaborn as sns
%matplotlib inline
sns.distplot(array)
```

<seaborn.axisgrid.FacetGrid at 0x78f3291c2710>



```
sns.distplot(array)
```


```
sns.distplot(array)
<Axes: ylabel='Density'>
```



```
new_array=array[(array>lr) & (array<ur)]
new_array
array([27, 50, 44, 6, 58, 61, 23, 86, 67, 20, 75, 7, 79, 61, 90, 54])
```

```
sns.displot(new_array)
```


```
<seaborn.axisgrid.FacetGrid at 0x78f2e09bb580>
```




```

lr1,ur1=outDetection(new_array)
lr1,ur1
(-38.5, 133.5)
final_array=new_array[(new_array>lr1) & (new_array<ur1)]
final_array
array([27, 50, 44, 6, 58, 61, 23, 86, 67, 20, 75, 7, 79, 61, 90, 54])
sns.distplot(final_array)

```



Lab experiments
Roll no: 230701017
Name: Ajay Srinivas R
Class: CSE-A II
Subject: Fundamentals of data science (CS23334)

Experiment: 05

```

import numpy as np

import pandas as pd

df=pd.read_csv("Hotel_Dataset.csv")

df

```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
9	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
10	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

df.duplicated()

```
0    False
1    False
2    False
3    False
4    False
5    False
6    False
7    False
8    False
9     True
10   False
dtype: bool
```

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11 entries, 0 to 10
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            11 non-null    int64
1   Age_Group              11 non-null    object
2   Rating(1-5)           11 non-null    int64
3   Hotel                  11 non-null    object
4   FoodPreference         11 non-null    object
5   Bill                   11 non-null    int64
6   NoOfPax                11 non-null    int64
7   EstimatedSalary        11 non-null    int64
8   Age_Group.1            11 non-null    object
dtypes: int64(5), object(4)
memory usage: 924.0+ bytes
```

df.drop_duplicates(inplace=True)

df

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
10	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

```
len(df)
```

```
10
```

```
index=np.array(list(range(0,len(df))))
```

```
df.set_index(index,inplace=True)
```

```
index
```

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
df
```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
9	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

```
df.drop(['Age_Group.1'],axis=1,inplace=True)
```

```
df
```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1	20-25	4	Ibis	veg	1300	2	40000
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000
2	3	25-30	6	RedFox	Veg	1322	2	30000
3	4	20-25	-1	LemonTree	Veg	1234	2	120000
4	5	35+	3	Ibis	Vegetarian	989	2	45000
5	6	35+	3	Ibys	Non-Veg	1909	2	122220
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122
7	8	20-25	7	LemonTree	Veg	2999	-10	345673
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999
9	10	30-35	5	RedFox	non-Veg	-6755	4	87777

df.CustomerID.loc[df.CustomerID<0]=np.nan

df.Bill.loc[df.Bill<0]=np.nan

df.EstimatedSalary.loc[df.EstimatedSalary<0]=np.nan

df

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4.0	Ibis	veg	1300.0	2	40000.0
1	2.0	30-35	5.0	LemonTree	Non-Veg	2000.0	3	59000.0
2	3.0	25-30	NaN	RedFox	Veg	1322.0	2	30000.0
3	4.0	20-25	NaN	LemonTree	Veg	1234.0	2	120000.0
4	5.0	35+	3.0	Ibis	Vegetarian	989.0	2	45000.0
5	6.0	35+	3.0	Ibys	Non-Veg	1909.0	2	122220.0
6	7.0	35+	4.0	RedFox	Vegetarian	1000.0	-1	21122.0
7	8.0	20-25	NaN	LemonTree	Veg	2999.0	-10	345673.0
8	9.0	25-30	2.0	Ibis	Non-Veg	3456.0	3	NaN
9	10.0	30-35	5.0	RedFox	non-Veg	NaN	4	87777.0

df['NoOfPax'].loc[(df['NoOfPax']<1) | (df['NoOfPax']>20)]=np.nan

df

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4.0	Ibis	veg	1300.0	2.0	40000.0
1	2.0	30-35	5.0	LemonTree	Non-Veg	2000.0	3.0	59000.0
2	3.0	25-30	NaN	RedFox	Veg	1322.0	2.0	30000.0
3	4.0	20-25	NaN	LemonTree	Veg	1234.0	2.0	120000.0
4	5.0	35+	3.0	Ibis	Vegetarian	989.0	2.0	45000.0
5	6.0	35+	3.0	Ibys	Non-Veg	1909.0	2.0	122220.0
6	7.0	35+	4.0	RedFox	Vegetarian	1000.0	NaN	21122.0
7	8.0	20-25	NaN	LemonTree	Veg	2999.0	NaN	345673.0
8	9.0	25-30	2.0	Ibis	Non-Veg	3456.0	3.0	NaN
9	10.0	30-35	5.0	RedFox	non-Veg	NaN	4.0	87777.0

```
df.Age_Group.unique()
```

```
array(['20-25', '30-35', '25-30', '35+'], dtype=object)
```

```
df.Hotel.unique()
```

```
array(['Ibis', 'LemonTree', 'RedFox', 'Ibys'], dtype=object)
```

```
df.Hotel.replace(['Ibys'],'Ibis',inplace=True)
```

```
df.FoodPreference.unique
```

```
<bound method Series.unique of 0 veg
```

```
1 Non-Veg
```

```
2 Veg
```

```
3 Veg
```

```
4 Vegetarian
```

```
5 Non-Veg
```

```
6 Vegetarian
```

```
7 Veg
```

```
8 Non-Veg
```

```
9 non-Veg
```

```
Name: FoodPreference, dtype: object>
```

```
df.FoodPreference.replace(['Vegetarian','veg'],'Veg',inplace=True)
```

```
df.FoodPreference.replace(['non-Veg'],'Non-Veg',inplace=True)
```

```
df.EstimatedSalary.fillna(round(df.EstimatedSalary.mean()),inplace=True)
```

```
df.NoOfPax.fillna(round(df.NoOfPax.median()),inplace=True)
```

```
df['Rating(1-5)'].fillna(round(df['Rating(1-5)'].median()), inplace=True)
```

```
df.Bill.fillna(round(df.Bill.mean()),inplace=True)
```

```
df
```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4.0	Ibis	Veg	1300.0	2.0	40000.0
1	2.0	30-35	5.0	LemonTree	Non-Veg	2000.0	3.0	59000.0
2	3.0	25-30	4.0	RedFox	Veg	1322.0	2.0	30000.0
3	4.0	20-25	4.0	LemonTree	Veg	1234.0	2.0	120000.0
4	5.0	35+	3.0	Ibis	Veg	989.0	2.0	45000.0
5	6.0	35+	3.0	Ibis	Non-Veg	1909.0	2.0	122220.0
6	7.0	35+	4.0	RedFox	Veg	1000.0	2.0	21122.0
7	8.0	20-25	4.0	LemonTree	Veg	2999.0	2.0	345673.0
8	9.0	25-30	2.0	Ibis	Non-Veg	3456.0	3.0	96755.0
9	10.0	30-35	5.0	RedFox	Non-Veg	1801.0	4.0	87777.0

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 06

```
import numpy as np
import pandas as pd
df=pd.read_csv('/content/pre-process_datasample.csv')
```

```
df
```

	Country	Age	Salary	Purchased
--	---------	-----	--------	-----------

0	France	44.0	72000.0	No
---	--------	------	---------	----

1 Spain 27.0 48000.0 Yes
 2 Germany 30.0 54000.0 No
 3 Spain 38.0 61000.0 No
 4 Germany 40.0 NaN Yes
 5 France 35.0 58000.0 Yes
 6 Spain NaN 52000.0 No
 7 France 48.0 79000.0 Yes
 8 NaN 50.0 83000.0 No
 9 France 37.0 67000.0 Yes

Next steps: `df.head()`

	Country	Age	Salary	Purchased	
	France	44.0	72000.0		0
No	1	Spain	27.0		
48000.0		Yes		2	
Germany		30.0			
54000.0	No	3	Spain		
38.0	61000.0	No	4		
Germany	40	0	NaN		
Yes					

```
df.Country.fillna(df.Country.mode()[0],inplace=True)
features=df.iloc[:, :-1].values
```

```
df.Country.fillna(df.Country.mode()[0],inplace=True)
```

```
label=df.iloc[:,-1].values
from sklearn.impute import SimpleImputer

age=SimpleImputer(strategy="mean",missing_values=np.nan)

Salary=SimpleImputer(strategy="mean",missing_values=np.nan)

age.fit(features[:,[1]])
```

▼ SimpleImputer ⓘ ⓘ
SimpleImputer()

```
Salary.fit(features[:,[2]])
```

▼ SimpleImputer ⓘ ⓘ
SimpleImputer()

```
SimpleImputer()
```

▼ SimpleImputer ⓘ ⓘ
SimpleImputer()

```
features[:,[1]]=age.transform(features[:,[1]])
```

```
features[:,[2]]=Salary.transform(features[:,[2]])
```

```
features
```

```
array([[ 'France', 44.0, 72000.0],
 [ 'Spain', 27.0, 48000.0],
 [ 'Germany', 30.0, 54000.0],
 [ 'Spain', 38.0, 61000.0],
 [ 'Germany', 40.0, 63777.77777777778],
 [ 'France', 35.0, 58000.0],
 [ 'Spain', 38.77777777777778, 52000.0],
 [ 'France', 48.0, 79000.0],
 [ 'France', 50.0, 83000.0],
 [ 'France', 37.0, 67000.0]], dtype=object)
```



```
from sklearn.preprocessing import OneHotEncoder
```

```
oh = OneHotEncoder(sparse_output=False)
```

```
Country=oh.fit_transform(features[:,[0]])
```

```
Country
```

```
array([[1., 0., 0.],  
       [0., 0., 1.],  
       [0., 1., 0.],  
       [0., 0., 1.],  
       [0., 1., 0.],  
       [1., 0., 0.],  
       [0., 0., 1.],  
       [1., 0., 0.],  
  
       [1., 0., 0.],  
       [1., 0., 0.]])
```

```
final_set=np.concatenate((Country,features[:,[1,2]]),axis=1)
```

```
final_set
```

```
array([[1.0, 0.0, 0.0, 44.0, 72000.0],  
       [0.0, 0.0, 1.0, 27.0, 48000.0],  
       [0.0, 1.0, 0.0, 30.0, 54000.0],  
       [0.0, 0.0, 1.0, 38.0, 61000.0],  
       [0.0, 1.0, 0.0, 40.0, 63777.77777777778],  
       [1.0, 0.0, 0.0, 35.0, 58000.0],  
       [0.0, 0.0, 1.0, 38.77777777777778, 52000.0],  
       [1.0, 0.0, 0.0, 48.0, 79000.0],  
       [1.0, 0.0, 0.0, 50.0, 83000.0],  
       [1.0, 0.0, 0.0, 37.0, 67000.0]], dtype=object)
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc=StandardScaler()
```

```
sc.fit(final_set)
```

```
feat_standard_scaler=sc.transform(final_set)
```

```
feat_standard_scaler
```

```
array([[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,  
        7.58874362e-01, 7.49473254e-01],
```

```

[-1.00000000e+00, -5.00000000e-01, 1.52752523e+00,
-1.71150388e+00, -1.43817841e+00],
[-1.00000000e+00, 2.00000000e+00, -6.54653671e-01,
-1.27555478e+00, -8.91265492e-01],
[-1.00000000e+00, -5.00000000e-01, 1.52752523e+00,
-1.13023841e-01, -2.53200424e-01],
[-1.00000000e+00, 2.00000000e+00, -6.54653671e-01,
1.77608893e-01, 6.63219199e-16],
[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
-5.48972942e-01, -5.26656882e-01],
[-1.00000000e+00, -5.00000000e-01, 1.52752523e+00,
0.00000000e+00, -1.07356980e+00],
[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
1.34013983e+00, 1.38753832e+00],
[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
1.63077256e+00, 1.75214693e+00],
[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
-2.58340208e-01, 2.93712492e-01]])

```

```

from sklearn.preprocessing import MinMaxScaler
mms=MinMaxScaler(feature_range=(0,1))
mms.fit(final_set)
feat_minmax_scaler=mms.transform(final_set)
feat_minmax_scaler

```

```

array([[1. , 0. , 0. , 0.73913043, 0.68571429],
[0. , 0. , 1. , 0. , 0. ],
[0. , 1. , 0. , 0.13043478, 0.17142857],
[0. , 0. , 1. , 0.47826087, 0.37142857],
[0. , 1. , 0. , 0.56521739, 0.45079365],
[1. , 0. , 0. , 0.34782609, 0.28571429],
[0. , 0. , 1. , 0.51207729, 0.11428571],
[1. , 0. , 0. , 0.91304348, 0.88571429],
[1. , 0. , 0. , 1. , 1. ],
[1. , 0. , 0. , 0.43478261, 0.54285714]])

```

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 07

```
import numpy as np
```

```
import pandas as pd
df=pd.read_csv("/content/pre-process_datasample.csv")
df
```

	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes
2	Germany	30.0	54000.0	No
3	Spain	38.0	61000.0	No
4	Germany	40.0	NaN	Yes
5	France	35.0	58000.0	Yes
6	Spain	NaN	52000.0	No
7	France	48.0	79000.0	Yes
8	NaN	50.0	83000.0	No
9	France	37.0	67000.0	Yes

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
 # Column Non-Null Count Dtype
---  ---
 0 Country 9 non-null object
 1 Age 9 non-null float64
 2 Salary 9 non-null float64
 3 Purchased 10 non-null object
dtypes: float64(2), object(2)
memory usage: 448.0+ bytes
```

```
df.Country.mode()
```

Country
0 France

```
df.Country.mode()[0]
```

```
type(df.Country.mode())
```

```
df.Country.fillna(df.Country.mode()[0],inplace=True)

df.Age.fillna(df.Age.median(),inplace=True)

df.Salary.fillna(round(df.Salary.mean()),inplace=True)

df
```

	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes
2	Germany	30.0	54000.0	No
3	Spain	38.0	61000.0	No
4	Germany	40.0	63778.0	Yes
5	France	35.0	58000.0	Yes
6	Spain	38.0	52000.0	No
7	France	48.0	79000.0	Yes
8	France	50.0	83000.0	No
9	France	37.0	67000.0	Yes

```
pd.get_dummies(df.Country)
```

	France	Germany	Spain
0	True	False	False
1	False	False	True
2	False	True	False
3	False	False	True
4	False	True	False
5	True	False	False
6	False	False	True
7	True	False	False
8	True	False	False
9	True	False	False

```
updated_dataset=pd.concat([pd.get_dummies(df.Country),df.iloc[:,[1,2,3]]],axis=1)
```

```
updated_dataset
```

France Germany Spain Age Salary Purchased

0 True False False 44.0 72000.0 No
1 False False True 27.0 48000.0 Yes
2 False True False 30.0 54000.0 No
3 False False True 38.0 61000.0 No
4 False True False 40.0 63778.0 Yes
5 True False False 35.0 58000.0 Yes
6 False False True 38.0 52000.0 No
7 True False False 48.0 79000.0 Yes
8 True False False 50.0 83000.0 No
9 True False False 37 0 67000 0 Yes

df.info()

updated_dataset.Purchased.replace(['No', 'Yes'],[0,1],inplace=True)

updated_dataset

France Germany Spain Age Salary Purchased

0 True False False 44.0 72000.0 0
1 False False True 27.0 48000.0 1
2 False True False 30.0 54000.0 0
3 False False True 38.0 61000.0 0
4 False True False 40.0 63778.0 1
5 True False False 35.0 58000.0 1
6 False False True 38.0 52000.0 0
7 True False False 48.0 79000.0 1
8 True False False 50.0 83000.0 0
9 True False False 37 0 67000 0 1

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 08


```
import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
tips=sns.load_dataset('tips')
```


```
tips.head()
```

```
   total_bill  tip  sex  smoker  day  time  size
0    16.99   1.01 Female    No  Sun  Dinner    2
1    10.34   1.66  Male    No  Sun  Dinner    3
2    21.01   3.50  Male    No  Sun  Dinner    3
3    23.68   3.31  Male    No  Sun  Dinner    2
4    24.59   3.61 Female    No  Sun  Dinner    4
sns.displot(tips.total_bill,kde=True)
```


```
<seaborn.axisgrid.FacetGrid at 0x79bb4c7ea680>
```




```
sns.displot(tips.total_bill,kde=False)
```




```
sns.jointplot(x=tips.tip,y=tips.total_bill)  
<seaborn.axisgrid.JointGrid at 0x79bb08fc96c0>
```



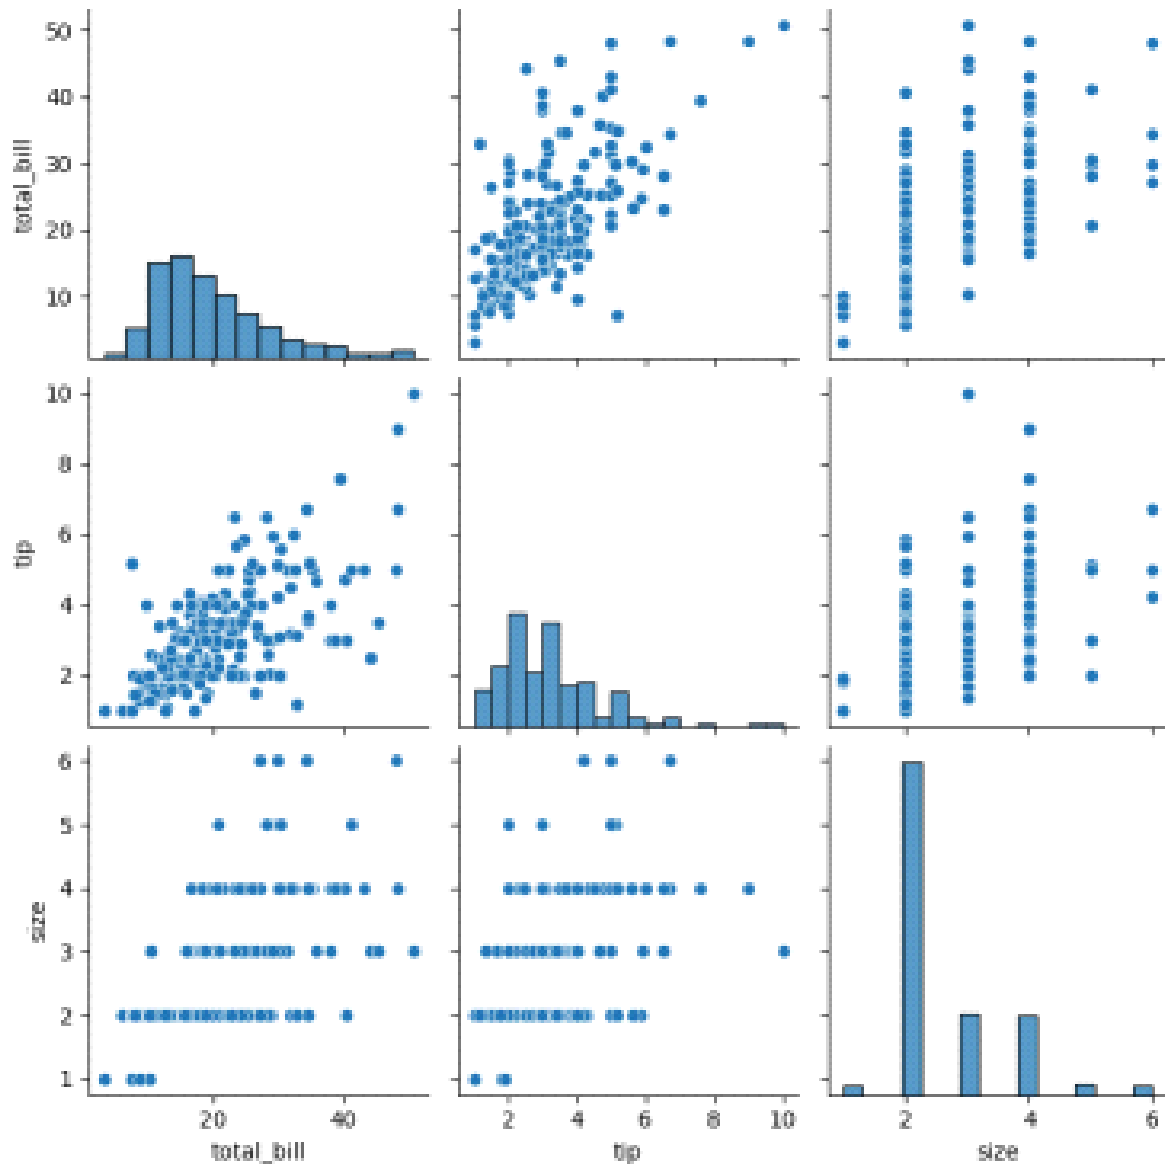
```
sns.jointplot(x=tips.tip,y=tips.total_bill,kind="reg")
```

```
sns.jointplot(x=tips.tip,y=tips.total_bill,kind="hex")  
<seaborn.axisgrid.JointGrid at 0x79bb088f4730>
```



```
sns.pairplot(tips)
```



```
tips.time.value_counts()
count
```

```
time
```

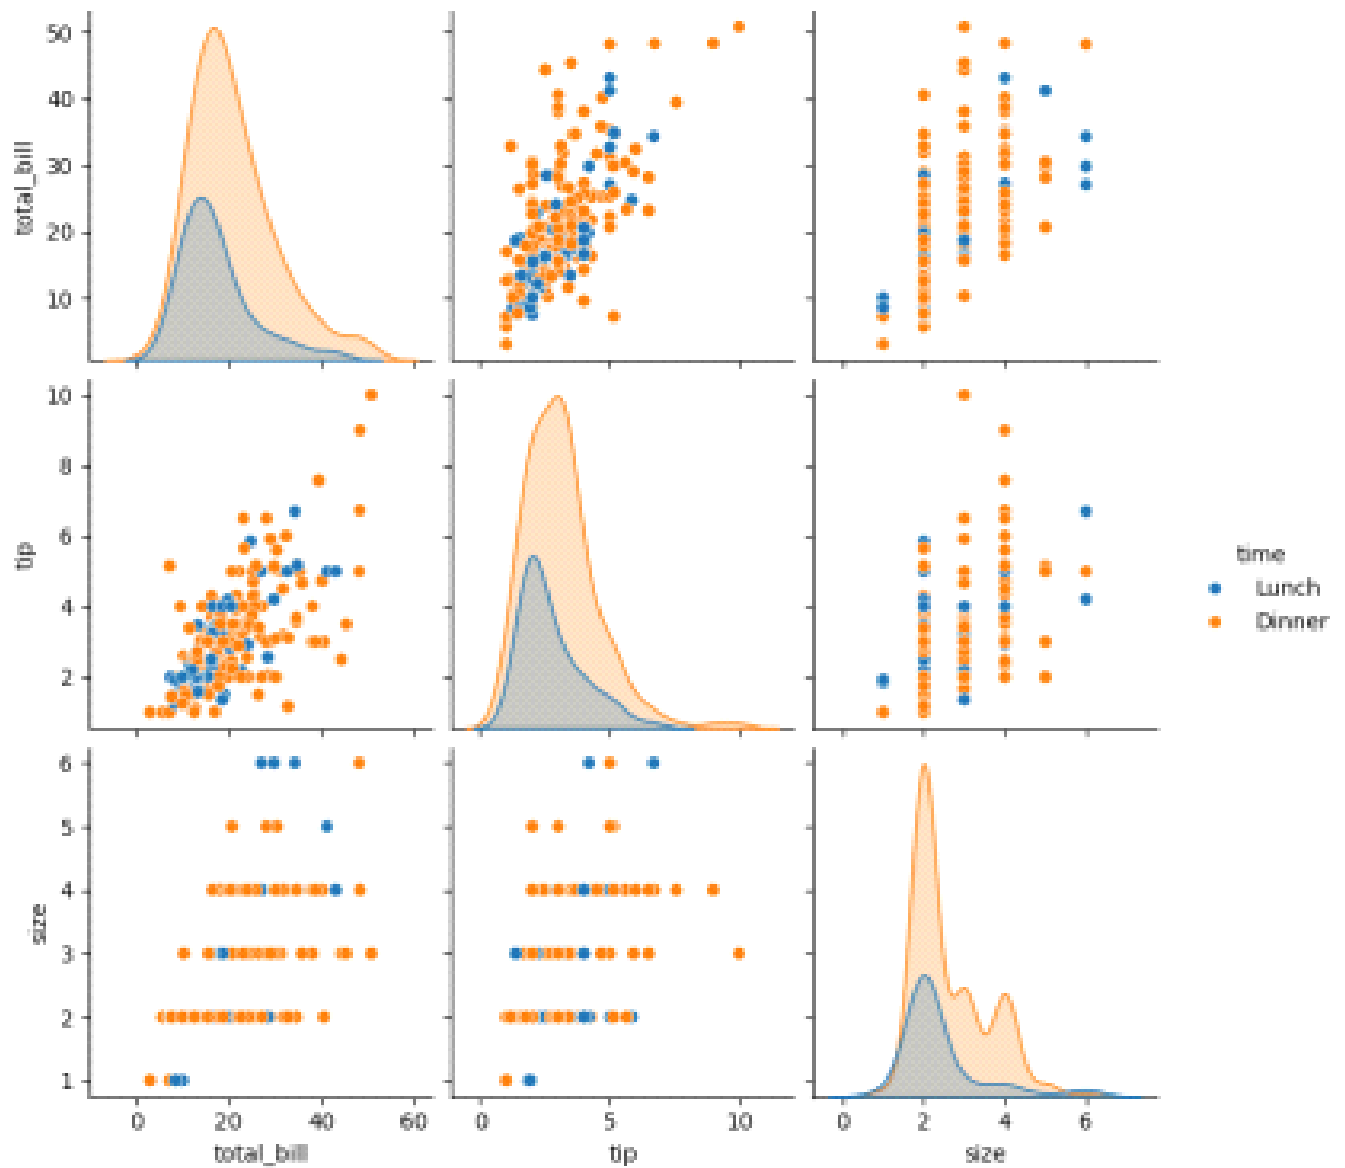
```
Dinner 176
```

```
Lunch 68
```

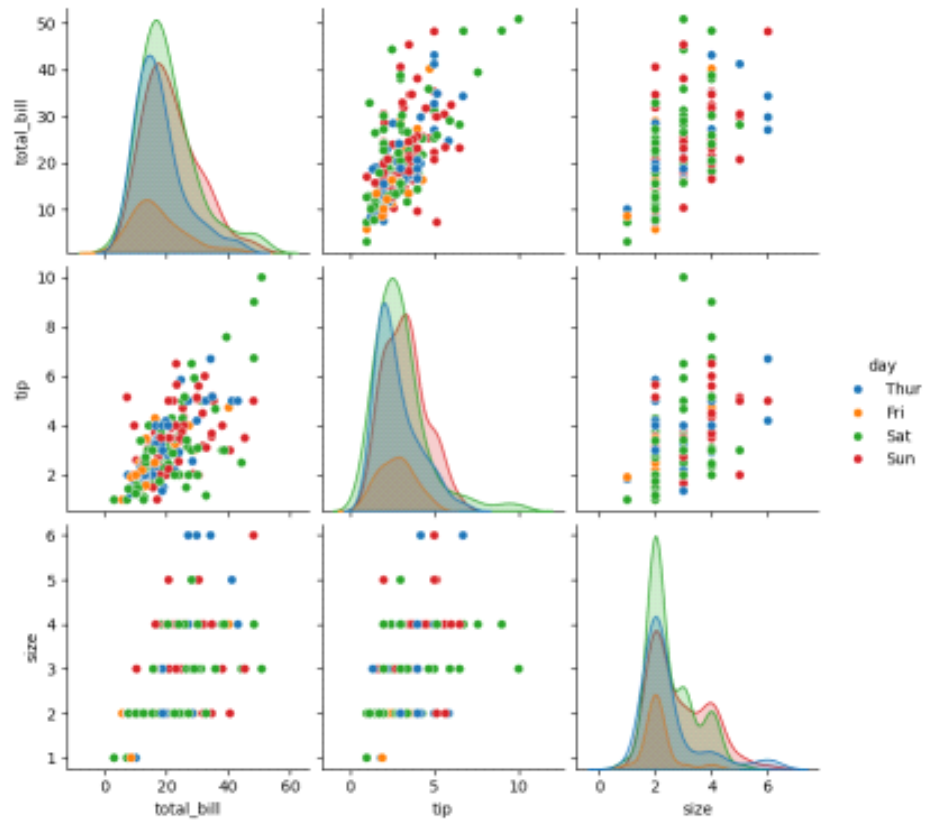
```
dtype: int64
```

```
sns.pairplot(tips,hue='time')
```

```
<seaborn.axisgrid.PairGrid at 0x79bb088f4670>
```




```
sns.pairplot(tips, hue='day')
```




```
sns.heatmap(tips.corr(numeric_only=True),annot=True)
```

<Axes: >




```
sns.boxplot(tips.total_bill)
```

<Axes: ylabel='total_bill'>




```
sns.boxplot(tips.tip)
```

<Axes: ylabel='tip'>





```
sns.countplot(tips.day)
```

<Axes: xlabel='count', ylabel='day'>




```
sns.countplot(tips.sex)
```





```
tips.sex.value_counts().plot(kind='pie')  
<Axes: ylabel='count'>
```

```
tips.sex.value_counts().plot(kind='bar')  
<Axes: xlabel='sex'>
```

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 08

```
# Column Non-Null Count Dtype ---
YearsExperience 30 non-null float64 1 Salary 30 non-null int64 dtypes:
float64(1), int64(1)
memory usage: 612.0 bytes
```

```
df.dropna(inplace=True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'> RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
# Column Non-Null Count Dtype ---
YearsExperience 30 non-null float64 1 Salary 30 non-null int64 dtypes:
float64(1), int64(1)
memory usage: 612.0 bytes
```

```
df.describe()
```

```
Out[5]: YearsExperience Salary
count 30.000000 30.000000
mean 5.313333 76003.000000
std 2.837888 27414.429785
min 1.100000 37731.000000
25% 3.200000 56720.750000
50% 4.700000 65237.000000
75% 7.700000 100544.750000
max 10.500000 122391.000000
```

In [6]:

```
features=df.iloc[:,[0]].values
label=df.iloc[:,[1]].values
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,random_st
```

```
from sklearn.linear_model import LinearRegression
model=LinearRegression()
model.fit(x_train,y_train)
```

```
Out[20]: ▾ LinearRegression
LinearRegression()
```

In [21]:

```
model.score(x_train,y_train)
```

```
Out[21]: 0.9603182547438908
```

In [23]:

```
model.score(x_test,y_test)
```

```
Out[23]: 0.9184170849214232
```

In [24]:

```
model.coef_
```

```
Out[24]: array([[9281.30847068]])
```

```
In [25]:  
model.intercept_
```

```
Out[25]: array([27166.73682891])
```

```
In [26]:  
import pickle  
pickle.dump(model,open('SalaryPred.model','wb'))
```

```
model=pickle.load(open('SalaryPred.model','rb')) yr_of_exp=float(input("Enter  
Years of Experience: "))  
yr_of_exp_NP=np.array([[yr_of_exp]])  
Salary=model.predict(yr_of_exp_NP)
```

```
Enter Years of Experience: 44
```

```
print("Estimated Salary for {} years of experience is {}: "
```

```
.format(yr_of_exp,Salary) Estimated Salary for 44.0 years of experience is
```

```
[[435544.30953887]]:
```

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 09

```
import numpy as np  
import pandas as pd
```

```
df=pd.read_csv('Iris.csv')  
df.info()
```

```
df.variety.value_counts()
```

```
Out[3]: Setosa 50
        Versicolor 50
        Virginica 50
        Name: variety, dtype: int64
```

```
In [4]:
df.head()
```

```
Out[4]: sepal.length sepal.width petal.length petal.width variety 0 5.1
        3.5 1.4 0.2 Setosa 1 4.9 3.0 1.4 0.2 Setosa 2
        4.7 3.2 1.3 0.2 Setosa 3 4.6 3.1 1.5 0.2
        Setosa 4 5.0 3.6 1.4 0.2 Setosa
```

```
In [5]: In [6]: In [8]:
```

```
features=df.iloc[:, :-1].values
label=df.iloc[:, 4].values
```

```
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
```

```
xtrain,xtest,ytrain,ytest=train_test_split(features,label,test_size=.2,random_state=0)
model_KNN=KNeighborsClassifier(n_neighbors=5)
model_KNN.fit(xtrain,ytrain)
```

```
Out[8]: KNeighborsClassifier()
        print(model_KNN.score(xtrain,ytrain))
        print(model_KNN.score(xtest,ytest))

0.9583333333333334
1.0
```

```
from sklearn.metrics import confusion_matrix
confusion_matrix(label,model_KNN.predict(features))
```

```
Out[10]: array([[50, 0, 0],
                [ 0, 47, 3],
                [ 0, 2, 48]], dtype=int64)
from sklearn.metrics import classification_report
print(classification_report(label,model_KNN.predict(features)))

precision recall f1-score support
```

```
Setosa 1.00 1.00 1.00 50 Versicolor 0.96 0.94 0.95 50 Virginica 0.94 0.96
0.95 50
```

```
accuracy 0.97 150 macro avg 0.97 0.97 0.97 150 weighted avg 0.97 0.97 0.97
150
```

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 10

In [1]:

```
import numpy as np
```

```
import pandas as pd
```

```
df=pd.read_csv('Social_Network_Ads.csv') df
```

Out[1]:

User	ID	Gender	Age	EstimatedSalary	Purchased		0
------	----	--------	-----	-----------------	-----------	--	---

15624510	Male	19	19000	0	1	15810944
----------	------	----	-------	---	---	----------

Male	35	20000	0	2	15668575	Female
------	----	-------	---	---	----------	--------

26	43000	0	3	15603246	Female	27
----	-------	---	---	----------	--------	----

57000	0	4	15804002	Male	19	76000	0
-------	---	---	----------	------	----	-------	---

... ..

395	15691863	Female	46	41000	1	396
-----	----------	--------	----	-------	---	-----

15706071	Male	51	23000	1	397
----------	------	----	-------	---	-----

15654296	Female	50	20000	1	398
----------	--------	----	-------	---	-----

15755018	Male	36	33000	0	399
----------	------	----	-------	---	-----

15594041	Female	49	36000	1
----------	--------	----	-------	---

400 rows x 5 columns

In [2]:

```
df.head()
```

Out[2]:

User	ID	Gender	Age	EstimatedSalary	Purchased
------	----	--------	-----	-----------------	-----------

0	15624510	Male	19	19000	0
---	----------	------	----	-------	---

1	15810944	Male	35	20000	0
---	----------	------	----	-------	---

2 15668575 Female 26 43000 0
3 15603246 Female 27 57000 0
4 15804002 Male 19 76000 0

In [4]:

```
features=df.iloc[:,[2,3]].values label=df.iloc[:,4].values features
```

```
Out[4]: array([[ 19,
                19000],
               [ 35,
                20000],
               [ 26,
                43000],
               [ 27,
                57000],
               [ 19,
                76000],
               [ 27,
                58000],
               [ 27,
                84000],
               [ 32,
                150000],
               [ 25,
                33000],
               [ 35,
                65000],
               [ 26,
                80000],
               [ 26,
                52000],
               [ 20,
                86000],
               [ 32,
                18000],
               [ 18,
                82000],
               [ 29,
                80000],
               [ 47,
                25000],
               [ 45,
                26000],
               [ 46,
                28000],
               [ 48,
                29000]
```

In [5]:

label

```
Out[5]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1,
                1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
                0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0,
                1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0,
                1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,
                1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1,
```

```

1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1,
1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1,
1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1,
1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1,
0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0,
1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1,
1, 0, 1], dtype=int64)

```

In [6]:

```

from sklearn.model_selection import train_test_split from
sklearn.linear_model import LogisticRegression

```

```

for i in range(1,401):
    x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.
model=LogisticRegression()
    model.fit(x_train,y_train)
    train_score=model.score(x_train,y_train)
    test_score=model.score(x_test,y_test)
    if test_score>train_score:
        print("Test {} Train{} Random State {}".format(test_score,train_score,i)

```

```

Test 0.6875 Train0.63125 Random State 3
Test 0.7375 Train0.61875 Random State 4
Test 0.6625 Train0.6375 Random State 5
Test 0.65 Train0.640625 Random State 6
Test 0.675 Train0.634375 Random State 7
Test 0.675 Train0.634375 Random State 8
Test 0.65 Train0.640625 Random State 10
Test 0.6625 Train0.6375 Random State 11
Test 0.7125 Train0.625 Random State 13
Test 0.675 Train0.634375 Random State 16
Test 0.7 Train0.628125 Random State 17
Test 0.7 Train0.628125 Random State 21
Test 0.65 Train0.640625 Random State 24
Test 0.6625 Train0.6375 Random State 25
Test 0.75 Train0.615625 Random State 26
Test 0.675 Train0.634375 Random State 27
Test 0.7 Train0.628125 Random State 28
Test 0.6875 Train0.63125 Random State 29
Test 0.6875 Train0.63125 Random State 31
T t 0 6625 T i 0 6375 R d St t 37

```

```

x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,
finalModel=LogisticRegression()
finalModel.fit(x_train,y_train)

```

Out[8]: LogisticRegression()

```

print(finalModel.score(x_train,y_train))

print(finalModel.score(x_test,y_test))

0.834375
0.9125

from sklearn.metrics import classification_report
print(classification_report(label,finalModel.predict(features)))

precision recall f1-score support

0 0.85 0.93 0.89 257  1 0.84 0.71 0.77 143

accuracy 0.85 400  macro avg 0.85 0.82 0.83 400 weighted avg 0.85 0.85 0.85
400

```

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 11

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

df=pd.read_csv('Mall_Customers.csv')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
# Column Non-Null Count Dtype ---
0 CustomerID 200 non-null int64 1 Gender 200 non-null object 2 Age 200 non-null int64 3
Annual Income (k$) 200 non-null int64 4 Spending Score (1-100) 200 non-null
int64 dtypes: int64(4), object(1)
memory usage: 7.9+ KB

```



```
df.head()
```

```
Out[4]: CustomerID Gender Age Annual Income (k$) Spending Score (1-100)
```

```
0 1 Male 19 15 39
```

```
1 2 Male 21 15 81
```

```
2 3 Female 20 16 6
```

```
3 4 Female 23 16 77
```

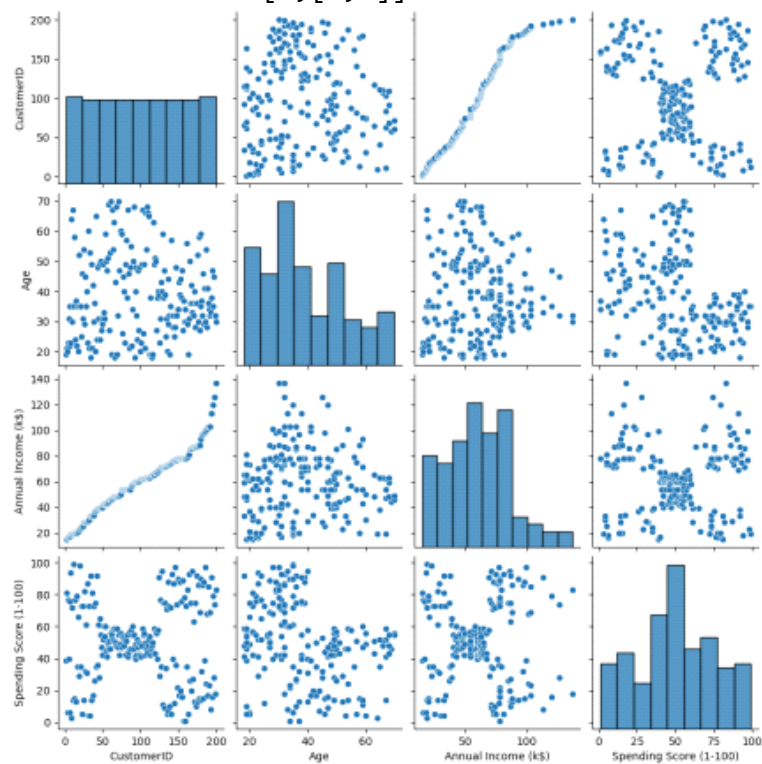
```
4 5 Female 31 17 40
```

```
sns.pairplot(df)
```

```
In [5]:
```

```
Out[5]: <seaborn.axisgrid.PairGrid at 0x170e8e47850>
```

```
features=df.iloc[:,[3,4]].values
```



```
In [6]:
```

```
In [7]:
from sklearn.cluster import KMeans
model=KMeans(n_clusters=5)
model.fit(features)
KMeans(n_clusters=5)
```

```
Out[7]: KMeans(n_clusters=5)
```

```
In [8]:
Final=df.iloc[:,[3,4]]
Final['label']=model.predict(features)
Final.head()
Final['label']=model.predict(features)
```

```
Out[8]: Annual Income (k$) Spending Score (1-100) label
```

```
0 15 39 4
```

```
1 15 81 2
```


```
2 16 6 4
```

```
3 16 77 2
```

```
4 17 40 4
```


```
In [9]: sns.set_style("whitegrid")
sns.FacetGrid(Final,hue="label",height=8) \
.map(plt.scatter,"Annual Income (k$)", "Spending Score (1-100)") \
```

```
.add_legend();  
plt.show()
```



```
In [10]: features_el=df.iloc[:,[2,3,4]].values
         from sklearn.cluster import KMeans
         wcss=[]
         for i in range(1,10):
             model=KMeans(n_clusters=i)
             model.fit(features_el)
             wcss.append(model.inertia_)
         plt.plot(range(1,10),wcss)
```

```
Out[10]: [<matplotlib.lines.Line2D at 0x170e99f3550>]
```



Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 12

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
# Step 1: Generate a population (e.g., normal distribution)
```

```
population_mean = 50
```

```
population_std = 10
```

```
population_size = 100000
```

```
population = np.random.normal(population_mean, population_std, population_size)
```

```
# Step 2: Random sampling
```

```
sample_sizes = [30, 50, 100] # different sample sizes to consider
```

```
num_samples = 1000 # number of samples for each sample size
```

```
sample_means = { }
```

```

for size in sample_sizes:
    sample_means[size] = []
    for _ in range(num_samples):
        sample = np.random.choice(population, size=size, replace=False)
        sample_means[size].append(np.mean(sample))


# Step 3: Plotting sampling distributions
plt.figure(figsize=(12, 8))

for i, size in enumerate(sample_sizes):
    plt.subplot(len(sample_sizes), 1, i+1)
    plt.hist(sample_means[size], bins=30, alpha=0.7, label=f'Sample Size {size}')
    plt.axvline(np.mean(population), color='red', linestyle='dashed', linewidth=1.5, label='Population Mean')
    plt.title(f'Sampling Distribution (Sample Size {size})')
    plt.xlabel('Sample Mean')
    plt.ylabel('Frequency')
    plt.legend()

plt.tight_layout()
plt.show()

```

OUTPUT:



Lab experiments

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 13

```
import numpy as np
```

```
import scipy.stats as stats
```

```
sample_data = np.array([152, 148, 151, 149, 147, 153, 150, 148, 152,  
                        149, 151, 150, 149, 152, 151, 148, 150, 152,  
                        149, 150, 148, 153, 151, 150, 149, 152,  
                        148, 151, 150, 153])
```

```
population_mean = 150
```

```
sample_mean = np.mean(sample_data)
```

```
sample_std = np.std(sample_data, ddof=1)
```

```
n = len(sample_data)
```

```
z_statistic = (sample_mean - population_mean) / (sample_std / np.sqrt(n))
```

```
p_value = 2 * (1 - stats.norm.cdf(np.abs(z_statistic)))
```

```
print(f"Sample Mean: {sample_mean:.2f}")
```

```
print(f"Z-Statistic: {z_statistic:.4f}")
```

```
print(f"P-Value: {p_value:.4f}")
```

```
alpha = 0.05
```

```
if p_value < alpha:
```

```
    print("Reject the null hypothesis: The average weight is significantly different from 150  
grams.")
```

```
else:
```

```
    print("Fail to reject the null hypothesis: There is no significant difference in average weight  
from 150 grams.")
```

OUTPUT:

Sample Mean: 150.20

Z-Statistic: 0.6406

P-Value: 0.5218

Fail to reject the null hypothesis: There is no significant difference in average weight from 150 grams.

Lab experiments

Roll no: 230701015

Name: Aishwarya A

Class: CSE-A II

Subject: Fundamentals of data science (CS2334)

Experiment: 13

```
import numpy as np
import scipy.stats as stats

# Set a random seed for reproducibility
np.random.seed(42)

# Generate hypothetical sample data (IQ scores)
sample_size = 25
sample_data = np.random.normal(loc=102, scale=15, size=sample_size) #
Mean IQ of 102, SD of 15

# Population mean under the null hypothesis
population_mean = 100

# Calculate sample statistics
sample_mean = np.mean(sample_data)
sample_std = np.std(sample_data, ddof=1) # Using sample standard
deviation

# Number of observations
n = len(sample_data)

# Calculate the T-statistic and p-value
t_statistic, p_value = stats.ttest_1samp(sample_data, population_mean)

# Print results
print(f"Sample Mean: {sample_mean:.2f}")
print(f"T-Statistic: {t_statistic:.4f}")
print(f"P-Value: {p_value:.4f}")
```



```
# Decision based on the significance level
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: The average IQ score is
significantly different from 100.")
else:
    print("Fail to reject the null hypothesis: There is no significant
difference in average IQ score from 100.")
```

OUTPUT:

```
Sample Mean: 99.55
T-Statistic: -0.1577
P-Value: 0.8760
Fail to reject the null hypothesis: There is no significant difference
in average IQ score from 100.
```

Lab experiments

Roll no: 230701017

Name: Ajay Srinivas R

Class: CSE-A II

Subject: Fundamentals of data science (CS23334)

Experiment: 14

```
import numpy as np
import scipy.stats as stats

# Set a random seed for reproducibility
np.random.seed(42)

# Generate hypothetical growth data for three treatments (A, B, C)
n_plants = 25

# Growth data (in cm) for Treatment A, B, and C
growth_A = np.random.normal(loc=10, scale=2, size=n_plants)
growth_B = np.random.normal(loc=12, scale=3, size=n_plants)
growth_C = np.random.normal(loc=15, scale=2.5, size=n_plants)

# Combine all data into one array
```

```

all_data = np.concatenate([growth_A, growth_B, growth_C])

# Treatment labels for each group
treatment_labels = ['A'] * n_plants + ['B'] * n_plants + ['C'] *
n_plants

# Perform one-way ANOVA
f_statistic, p_value = stats.f_oneway(growth_A, growth_B, growth_C)

# Print results
print("Treatment A Mean Growth:", np.mean(growth_A))
print("Treatment B Mean Growth:", np.mean(growth_B))
print("Treatment C Mean Growth:", np.mean(growth_C))
print()
print(f"F-Statistic: {f_statistic:.4f}")
print(f"P-Value: {p_value:.4f}")

# Decision based on the significance level
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant
difference in mean growth rates among the three treatments.")
else:
    print("Fail to reject the null hypothesis: There is no significant
difference in mean growth rates among the three treatments.")

# Additional: Post-hoc analysis (Tukey's HSD) if ANOVA is significant
if p_value < alpha:
    from statsmodels.stats.multicomp import pairwise_tukeyhsd

    tukey_results = pairwise_tukeyhsd(all_data, treatment_labels,
alpha=0.05)
    print("\nTukey's HSD Post-hoc Test:")
    print(tukey_results)

```

OUTPUT:

```

Treatment A Mean Growth: 9.672983882683818
Treatment B Mean Growth: 11.137680744437432
Treatment C Mean Growth: 15.265234904828972

```

F-Statistic: 36.1214

P-Value: 0.0000

Reject the null hypothesis: There is a significant difference in mean growth rates among the three treatments.